

```
# -*- coding: utf-8 -*-
"""9062.ipynb
```

Automatically generated by Colaboratory.

```
Original file is located at
    https://colab.research.google.com/drive/1aaFkTtKH9JeQIohAHG-mU5JvH4nFZqAS
"""
```

```
# Commented out IPython magic to ensure Python compatibility.
# %pip install -q tf-models-official==2.7.1 tensorflow-text==2.7.3
```

```
"""## Importing modules
- `numpy` is imported due to setting random seed for pandas & sklearn so as to
make it easier to reproduce results
- `pandas` is imported for reading/loading are dataset which is in a tabular form
(comma seperated value file format)
- `matplotlib` is imported for creating graphs to visualize the training of our
model. It's easier to see if it overfits or underfits
- We are using `sklearn` for dividing our dataset into 3 parts, ie; training,
validation and test sets
- We are importing modules for `tensorflow` for our Deep Learning model
  - main tensorflow module for making, compiling and training the model.
  - `tensorflow-hub` for Bert preprocessor and text encoder
  - `tensorflow-text` as a workaround for a OP call error in `tensorflow-hub`
possibly due to scoping issue or not being listed as a dependency
  - `tf-models-official` or official (as imported here) is used to get the AdamW
optimizer for stochastic gradient descent based training of out model
"""
```

```
# Commented out IPython magic to ensure Python compatibility.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
import tensorflow as tf
import tensorflow_text as tf_text
import tensorflow_hub as hub # for Bert Preprocessor and encoder
from official.nlp import optimization # for AdamW optimizer
```

```
tf.get_logger().setLevel('ERROR')
# %matplotlib inline
```

```
"""### Random Seed
Setting a seed manually here to keep things reproducible and consistent
"""
```

```

seed = 13
tf.random.set_seed(seed)
np.random.seed(seed)

"""### Dataset handling
Read the file and load it as a `pandas.DataFrame` for quick processing
"""

# load the dataset as a DataFrame
df = pd.read_csv("./deceptive-opinion.csv")

# Display 5 rows of tabular data, first 5 indices by default
df.head()

"""Describe gives some basic overview of the dataset by category/column/attribute,
such as
- `count`/size (total number of rows by each category, gives a heads up incase
there's missing attribute data),
- distinction (number of `unique` values for each attribute),
- `freq`uency of the top most (or most frequent value in an attribute)
"""

df.describe()

"""Just to confirm there are no null values, doesn't work with text/object
attribute wherein, it could be `"null"`/`{}`/`[]` instead of `nil`"""

df.isnull().sum()

"""- We are dropping `hotel` and review `source` attributes so that our model
learns to classify purely based on text and doesn't become brittle wherein it
breaks if it encounters a new hotel name or source for the review fed to it.

- By that we mean it doesn't end up overfitting on our dataset and be completely
impractical for use.

- `polarity` is removed due it being data that isn't readily available in the text
but rather an add-on value that has been provided in the dataset. We could make
another classifier for polarity and then graft it's output with original text to
use in our model but that's an unnecessary complication.
"""

new_df = df.drop(['hotel'], axis=1)
new_df = new_df.drop(['source'], axis=1)
new_df = new_df.drop(['polarity'], axis=1)
new_df.head()

"""We are making the text `lower` case since the classification is going to be case
insensitive for ease, since people may capitalise things for various reasons other
than grammatical as well as do it by mistake as is common on most user content on

```

the internet.

We are also removing `URL`s from the reviews (if found) since they aren't relevant and could confuse the model as a feature for identifying as it as spam.

"""

```
def text_cleaning(text):
```

```
    import re
```

```
    import string
```

```
    ...
```

```
    Make text lowercase, remove text in square brackets, remove links, remove special characters
```

```
    and remove words containing numbers.
```

```
    ...
```

```
    text = text.lower()
```

```
    text = re.sub('https?://\S+|www\.\S+', '', text) # remove URLs
```

```
    return text
```

```
new_df['text'] = new_df['text'].apply(text_cleaning)
```

"""Here we are mapping the `class`/category for truthful & deceptive to numeric form so that it's easier to measure loss, accuracy & confidence of our model."""

```
def className(cName):
```

```
    if cName == 'truthful':
```

```
        return 0
```

```
    else:
```

```
        return 1
```

```
new_df['deceptive'] = new_df['deceptive'].apply(className)
```

"""Taking advantage of jupyter notebook's output truncating feature we are displaying first 5 and last 5 rows/tuples of data after our processing of the dataset. This was the last of dataset preprocessing."""

```
new_df
```

"""- we are converting our tabular dataset from `pandas.DataFrame` to `numpy.array` for feeding our model & division by test, validation, training

- we also separate `classname` and `text` data as `y` and `x` `numpy.array` variables respectively from the tabular form it originally was

"""

```
x = new_df.copy(deep=True)
```

```
y = x.pop('deceptive').to_numpy()
```

```
x = x['text'].to_numpy()
```

```
x # text data array
```

y # class array: classes for each text, both are indexed exactly so x[0] text corresponds with y[0]

```
""" ## Train, Test, Validation splits
- Test set size: 160
- Validation set size: 144
- Train set size: 1296
"""
```

```
x_train, x_test, y_train, y_test = train_test_split(
    x, y, random_state=seed, test_size=0.1, stratify=y) # 1440, 160
x_train, x_val, y_train, y_val = train_test_split(
    x_train, y_train, random_state=seed, test_size=0.1, stratify=y_train) # 1296,
144
```

```
print(type(x_train), type(y_train))
```

```
training_set = pd.DataFrame({'text':x_train, 'class_label':y_train})
training_set.to_csv("training_set.csv", index=False)
validation_set = pd.DataFrame({'text':x_val, 'class_label':y_val})
validation_set.to_csv("validation_set.csv", index=False)
test_set = pd.DataFrame({'text':x_test, 'class_label':y_test})
test_set.to_csv("test_set.csv", index=False)
```

```
"""## Bert Encoder
We are using `small_Bert` with 4 hidden layers of 512 pooled output with 8
Attention heads
- `L` = number of hidden layers, valid values being one of `[2, 4, 6, 8, 10, 12]`
- `H` = pooled output size, valid values being one of `[128, 512, 768]`
- `A` = number of Attention head, valid values being one of `[2, 4, 8, 12]`
"""
```

```
bert_model_name = 'small_bert/bert_en_uncased_L-4_H-512_A-8'
```

```
map_name_to_handle = {
    'bert_en_uncased_L-12_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_uncased_L-12_H-768_A-12/3',
    'bert_en_cased_L-12_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_cased_L-12_H-768_A-12/3',
    'bert_multi_cased_L-12_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_multi_cased_L-12_H-768_A-12/3',
    'small_bert/bert_en_uncased_L-2_H-128_A-2':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-2_H-128_A-2/1',
    'small_bert/bert_en_uncased_L-2_H-256_A-4':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-2_H-256_A-4/1',
    'small_bert/bert_en_uncased_L-2_H-512_A-8':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-2_H-512_A-8/1',
    'small_bert/bert_en_uncased_L-2_H-768_A-12':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-2_H-768_A-12/1',
    'small_bert/bert_en_uncased_L-4_H-128_A-2':
```

'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-128_A-2/1',
'small_bert/bert_en_uncased_L-4_H-256_A-4':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-256_A-4/1',
'small_bert/bert_en_uncased_L-4_H-512_A-8':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-512_A-8/1',
'small_bert/bert_en_uncased_L-4_H-768_A-12':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-768_A-12/1',
'small_bert/bert_en_uncased_L-6_H-128_A-2':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-6_H-128_A-2/1',
'small_bert/bert_en_uncased_L-6_H-256_A-4':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-6_H-256_A-4/1',
'small_bert/bert_en_uncased_L-6_H-512_A-8':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-6_H-512_A-8/1',
'small_bert/bert_en_uncased_L-6_H-768_A-12':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-6_H-768_A-12/1',
'small_bert/bert_en_uncased_L-8_H-128_A-2':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-8_H-128_A-2/1',
'small_bert/bert_en_uncased_L-8_H-256_A-4':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-8_H-256_A-4/1',
'small_bert/bert_en_uncased_L-8_H-512_A-8':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-8_H-512_A-8/1',
'small_bert/bert_en_uncased_L-8_H-768_A-12':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-8_H-768_A-12/1',
'small_bert/bert_en_uncased_L-10_H-128_A-2':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-10_H-128_A-2/1',
'small_bert/bert_en_uncased_L-10_H-256_A-4':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-10_H-256_A-4/1',
'small_bert/bert_en_uncased_L-10_H-512_A-8':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-10_H-512_A-8/1',
'small_bert/bert_en_uncased_L-10_H-768_A-12':

'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-10_H-768_A-12/1',
'small_bert/bert_en_uncased_L-12_H-128_A-2':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-12_H-128_A-2/1',
'small_bert/bert_en_uncased_L-12_H-256_A-4':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-12_H-256_A-4/1',
'small_bert/bert_en_uncased_L-12_H-512_A-8':
'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-12_H-512_A-8/1',
'small_bert/bert_en_uncased_L-12_H-768_A-12':

'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-12_H-768_A-12/1',
'albert_en_base':
'https://tfhub.dev/tensorflow/albert_en_base/2',
'electra_small':
'https://tfhub.dev/google/electra_small/2',
'electra_base':
'https://tfhub.dev/google/electra_base/2',
'experts_pubmed':
'https://tfhub.dev/google/experts/bert/pubmed/2',
'experts_wiki_books':

```

        'https://tfhub.dev/google/experts/bert/wiki_books/2',
        'talking-heads_base':
            'https://tfhub.dev/tensorflow/talkheads_ggelu_bert_en_base/1',
    }

```

```

map_model_to_preprocess = {
    'bert_en_uncased_L-12_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'bert_en_cased_L-12_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_cased_preprocess/3',
    'small_bert/bert_en_uncased_L-2_H-128_A-2':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-2_H-256_A-4':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-2_H-512_A-8':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-2_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-4_H-128_A-2':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-4_H-256_A-4':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-4_H-512_A-8':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-4_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-6_H-128_A-2':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-6_H-256_A-4':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-6_H-512_A-8':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-6_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-8_H-128_A-2':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-8_H-256_A-4':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-8_H-512_A-8':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-8_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-10_H-128_A-2':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-10_H-256_A-4':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-10_H-512_A-8':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-10_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
}

```

```

'small_bert/bert_en_uncased_L-12_H-128_A-2':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-12_H-256_A-4':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-12_H-512_A-8':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-12_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'bert_multi_cased_L-12_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_multi_cased_preprocess/3',
'albert_en_base':
    'https://tfhub.dev/tensorflow/albert_en_preprocess/3',
'electra_small':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'electra_base':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'experts_pubmed':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'experts_wiki_books':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'talking-heads_base':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
}

```

```

tfhub_handle_encoder = map_name_to_handle[bert_model_name]
tfhub_handle_preprocess = map_model_to_preprocess[bert_model_name]

```

```

print(f'BERT model selected          : {tfhub_handle_encoder}')
print(f'Preprocess model auto-selected: {tfhub_handle_preprocess}')

```

""""Here we are just looking at what and how `Bert` model we selected in prior cells are handling a test input""""

```

bert_preprocess_model = hub.KerasLayer(tfhub_handle_preprocess)
text_test = ['this is such an amazing hotel, got free breakfast!']
text_preprocessed = bert_preprocess_model(text_test)

```

```

print(f'Keys          : {list(text_preprocessed.keys())}')
print(f'Shape         : {text_preprocessed["input_word_ids"].shape}')
print(f'Word Ids        : {text_preprocessed["input_word_ids"][0, :16]}')
print(f'Input Mask       : {text_preprocessed["input_mask"][0, :16]}')
print(f'Type Ids         : {text_preprocessed["input_type_ids"][0, :16]}')

```

```

bert_model = hub.KerasLayer(tfhub_handle_encoder)

```

```

bert_results = bert_model(text_preprocessed)

```

```

print(f'Loaded BERT: {tfhub_handle_encoder}')
print(f'Pooled Outputs Shape:{bert_results["pooled_output"].shape}')
print(f'Pooled Outputs Values:{bert_results["pooled_output"][0, :12]}')

```

```
print(f'Sequence Outputs Shape:{bert_results["sequence_output"].shape}')
print(f'Sequence Outputs Values:{bert_results["sequence_output"][0, :12]}')
```

```
"""## Model building
```

This is our classifier model composed of `Bert` Encoder (includes it's own preprocessing as well) and 1 Dimensional Convolution layers

```
(`tf.keras.layers.Conv1D`).
```

- We normalize the `pooled_output` from `Bert` before feeding it into the classifier

- It is then reshaped from `(1, 512)` to `(16, 32)` for convolution layers

- kernels(layer parameters) are initialised with

```
`tf.keras.initializers.GlorotUniform` for use with `relu` (rectified ELU)
```

activation

- `tf.keras.layers.Dropout` for regularization.

- `tf.keras.layers.MaxPool` for using only the most pronounced features of encoded text

- `tf.keras.layers.Flatten` is used to flattening the output from prior layers for the final perceptron

- `tf.keras.layers.Dense` is the decision perceptron (Also referred to as `FullyConnected`) with sigmoid activation to get an output between `(0, 1)`

```
"""
```

```
def build_classifier_model():
```

```
    text_input = tf.keras.layers.Input(shape=(), dtype=tf.string, name='text')
```

```
    preprocessing_layer = hub.KerasLayer(
```

```
        tfhub_handle_preprocess, name='preprocessing')
```

```
    encoder_inputs = preprocessing_layer(text_input) # [1, 128]
```

```
    encoder = hub.KerasLayer(tfhub_handle_encoder,
```

```
        trainable=True, name='BERT_encoder')
```

```
    outputs = encoder(encoder_inputs)
```

```
    net = outputs['pooled_output'] # [1, 512]
```

```
    initializer = tf.keras.initializers.GlorotUniform(seed=seed)
```

```
    net = tf.keras.layers.LayerNormalization(axis=1)(net)
```

```
    net = tf.keras.layers.Reshape((16, 32))(net)
```

```
    net = tf.keras.layers.Dropout(0.25)(net)
```

```
    net = tf.keras.layers.Conv1D(
```

```
        filters=128, kernel_size=9, kernel_initializer=initializer, padding='same',
```

```
activation='relu')(net)
```

```
    net = tf.keras.layers.Dropout(0.25)(net)
```

```
    net = tf.keras.layers.MaxPool1D(pool_size=2)(net)
```

```
    net = tf.keras.layers.Conv1D(
```

```
        filters=128, kernel_size=7, kernel_initializer=initializer, padding='same',
```

```
activation='relu')(net)
```

```
    net = tf.keras.layers.Dropout(0.25)(net)
```

```
    net = tf.keras.layers.MaxPool1D(pool_size=2)(net)
```

```
    net = tf.keras.layers.Flatten()(net)
```

```
    net = tf.keras.layers.Dense(
```

```
        1, activation='sigmoid', name='classifier')(net)
```

```
    return tf.keras.Model(text_input, net)
```



```

# Test code to see if model is able to process simple text string or not as
intended
classifier_model = build_classifier_model()
bert_raw_result = classifier_model(tf.constant(text_test))
# This result here isn't accurate as model hasn't been trained yet
# It's just to check input/output of the model
print(bert_raw_result)

tf.keras.utils.plot_model(classifier_model, show_shapes=True, show_dtype=True,
                           show_layer_names=True, expand_nested=True, dpi=128,
                           show_layer_activations=True
                           )

""" - We are using `tf.keras.losses.BinaryCrossentropy` for Binary cross-entropy
loss for our binary classifier
- We are using EarlyStopping callback to prevent overfitting.
- `tf.metrics.BinaryAccuracy` is just a metric to evaluate model accuracy, binary
means it expects output between `(0, 1)`
"""

loss = tf.keras.losses.BinaryCrossentropy(from_logits=False)
metrics = [tf.metrics.BinaryAccuracy(), tf.keras.metrics.FalseNegatives(),
tf.keras.metrics.FalsePositives(), tf.keras.metrics.TrueNegatives(),
tf.keras.metrics.TruePositives()]
# minimum delta is an arbitrary 3e-3, consider 2e-3 or higher values as well
callbacks = [
    tf.keras.callbacks.EarlyStopping(
        monitor='val_loss', min_delta=1e-3, patience=3, verbose=0, mode='min',
        restore_best_weights=True),
    # tf.keras.callbacks.ModelCheckpoint("weights.hdf5", monitor='val_loss',
    save_best_only=True, save_freq='epoch'),
]

# Increased Epoch to 100 because EarlyStopping Callback will stop training once
improvements stop
epochs = 100

# Here we are initialising the `AdamW` optimizer with hyperparameters
steps_per_epoch = x_train.size
num_train_steps = steps_per_epoch * epochs
num_warmup_steps = int(0.1 * num_train_steps)

init_lr = 3e-5 # initial learning rate
optimizer = optimization.create_optimizer(
    init_lr=init_lr, num_train_steps=num_train_steps,
    num_warmup_steps=num_warmup_steps, optimizer_type='adamw')

# Compiles the model with the optimizer, loss function and metric to be used for
evaluating progress
classifier_model.compile(optimizer=optimizer, loss=loss, metrics=metrics)

```

```

"""## Model Training
This is the training step for our model
`history` variable stores the per `epoch` results for training and validation
dataset accuracy and loss
`step` = number of batches the the data is fed to the model for loss evaluation and
parameter optimization
`epoch` = 1 complete run of all available `steps`, whole training data has been fed
at the end of an epoch.
"""

print(f'Training model with {tfhub_handle_encoder}')
history = classifier_model.fit(x=x_train, y=y_train, validation_data=(
    x_val, y_val), epochs=epochs, callbacks=callbacks)

history_dict = history.history
print(history_dict.keys())

# a blind test for the model as it's never been evaluated against
# or trained with the test set for a test of how well it's learnt to classify

loss, accuracy, fn, fp, tn, tp = classifier_model.evaluate(x_test, y_test)

print(f'Loss: {loss:.2}')
print(f'Accuracy: {accuracy:.2}')
print(f'True Positive Rate: {tp / (fn + tp):.2}')
print(f'True Negative Rate: {tn / (fp + tn):.2}')
recall = tp / (fn + tp)
precision = tp / (fp + tp)
print(f'F1 Score: {(2 * precision * recall) / (precision + recall):.2}')
print(f'Positive Predictive Value: {tp / (fp + tp):.2}')
print(f'Negative Predictive Value: {tn / (fn + tn):.2}')

"""## Training Graph
This graph plots the progress of the model's training based `accuracy` & `loss` for
both training and validation datasets per `epoch`
"""

acc = history_dict['binary_accuracy']
val_acc = history_dict['val_binary_accuracy']
loss = history_dict['loss']
val_loss = history_dict['val_loss']

epochs = range(1, len(acc) + 1)
fig = plt.figure(figsize=(10, 6))
fig.tight_layout()

plt.subplot(2, 1, 1)
# r is for "solid red line"
plt.plot(epochs, loss, 'r', label='Training loss')

```

```

# b is for "solid blue line"
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='upper right')

plt.subplot(2, 1, 2)
plt.plot(epochs, acc, 'r', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')

acc = np.array(history_dict['true_negatives']) /
(np.array(history_dict['false_positives']) +
np.array(history_dict['true_negatives']))
val_acc = np.array(history_dict['val_true_negatives']) /
(np.array(history_dict['val_false_positives']) +
np.array(history_dict['val_true_negatives']))
loss = np.array(history_dict['true_positives']) /
(np.array(history_dict['false_negatives']) +
np.array(history_dict['true_positives']))
val_loss = np.array(history_dict['val_true_positives']) /
(np.array(history_dict['val_false_negatives']) +
np.array(history_dict['val_true_positives']))

epochs = range(1, len(acc) + 1)
fig = plt.figure(figsize=(10, 6))
fig.tight_layout()

plt.subplot(2, 1, 1)
# r is for "solid red line"
plt.plot(epochs, loss, 'r', label='Training TNR')
# b is for "solid blue line"
plt.plot(epochs, val_loss, 'b', label='Validation TNR')
plt.title('Training and validation TNR')
plt.xlabel('Epochs')
plt.ylabel('True Negative Rate')
plt.legend(loc='lower right')

plt.subplot(2, 1, 2)
plt.plot(epochs, acc, 'r', label='Training TPR')
plt.plot(epochs, val_acc, 'b', label='Validation TPR')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')

```

```
from sklearn.metrics import ConfusionMatrixDisplay, RocCurveDisplay

y_pred = classifier_model.predict(x_test)
# print(type(y_pred), y_pred)
y_pred_discrete = [0 if y <= 0.5 else 1 for y in y_pred]
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_discrete)
plt.show()

RocCurveDisplay.from_predictions(y_test, y_pred)
plt.show()

classifier_model.save_weights("Weights.hdf5", overwrite=True)

classifier_model.save('Bert_Model', overwrite=True,
                     include_optimizer=True, save_format='tf')

!zip -r ./Bert_Model.zip ./Bert_Model
```